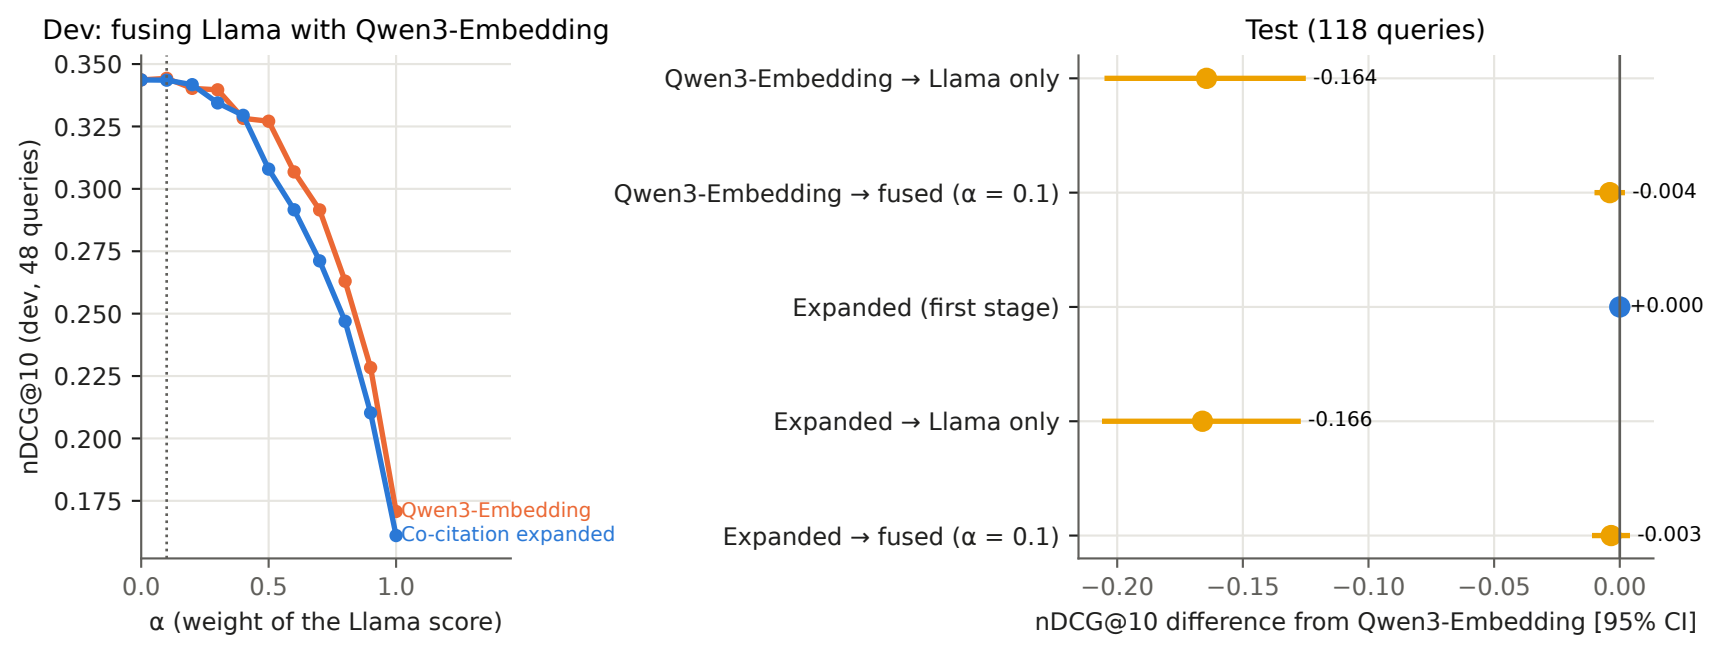
Dev: fusing Llama with Qwen3-Embedding
nDCG@10 (dev, 48 queries)
0.350
0.325
0.300
0.275
0.250
0.225
0.200
0.175
0.0
0.5
1.0
α (weight of the Llama score)
Qwen3-Embedding
Co-citation expanded
Test (118 queries)
Qwen3-Embedding → Llama only
-0.164
Qwen3-Embedding → fused (α = 0.1)
-0.004
Expanded (first stage)
+0.000
Expanded → Llama only
-0.166
Expanded → fused (α = 0.1)
-0.003
−0.20
−0.15
−0.10
−0.05
0.00
nDCG@10 difference from Qwen3-Embedding [95% CI]