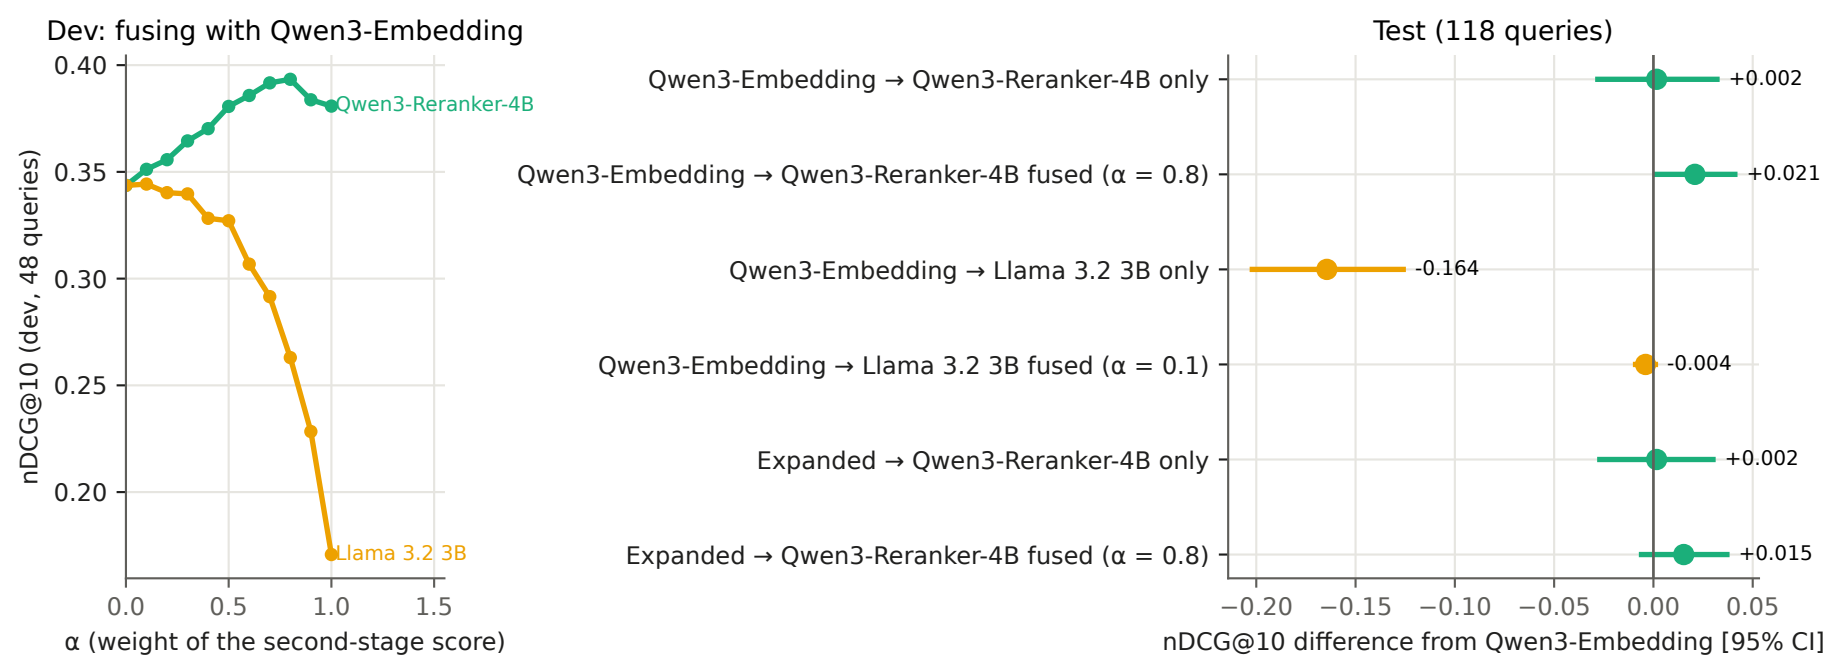
Dev: fusing with Qwen3-Embedding
nDCG@10 (dev, 48 queries)
0.40
0.35
0.30
0.25
0.20
Qwen3-Reranker-4B
Llama 3.2 3B
0.0
0.5
1.0
1.5
α (weight of the second-stage score)
Test (118 queries)
Qwen3-Embedding → Qwen3-Reranker-4B only
Qwen3-Embedding → Qwen3-Reranker-4B fused (α = 0.8)
Qwen3-Embedding → Llama 3.2 3B only
Qwen3-Embedding → Llama 3.2 3B fused (α = 0.1)
Expanded → Qwen3-Reranker-4B only
Expanded → Qwen3-Reranker-4B fused (α = 0.8)
+0.002
+0.021
-0.164
-0.004
+0.002
+0.015
−0.20
−0.15
−0.10
−0.05
0.00
0.05
nDCG@10 difference from Qwen3-Embedding [95% CI]